

National Information Ecosystem

Executive Brief

OS A01 A02

Leadership decision guide for responsible AI assisted statistical work

Decision required

Authorize a ninety day diagnostic of one priority statistical decision chain. The diagnostic must map actors, information assets, flows, user needs, quality risks and accountability before any operational AI deployment is approved.

Permit approved AI tools to accelerate bounded tasks such as evidence extraction, comparison, drafting and gap detection. Do not transfer conceptual, methodological, legal, domain or release authority to AI. A qualified statistician and the relevant subject matter specialist remain accountable for every consequential decision.

Control field	Executive position
Current authorization	Diagnostic and controlled screening only
AI role	Assist evidence work under explicit constraints and independent verification
Human authority	Statistician and subject matter specialist define criteria, validate meaning and approve disposition
Efficiency test	Time saved counts only when accepted quality, confidentiality and accountability are preserved

Executive conclusion

Official statistics are produced through a national information ecosystem rather than by one institution acting alone. Public agencies, administrative systems, research bodies, private data holders, technology providers, funders, oversight bodies and users jointly determine whether information is relevant, coherent, timely, interpretable and trusted.

The correct leadership response is not a general instruction to use AI. It is a controlled sequence: define the public decision; map actors, assets and flows; test fitness for use; resolve authority and safeguards; improve non AI dependencies; and screen only a bounded AI assisted task against a credible baseline.

What leaders should authorize now

Action	Leadership purpose	Minimum evidence
Scope one decision chain	Tie work to a real public decision and deadline.	Named user, product, population, geography, reference period and consequence of error.
Name accountable roles	Separate ownership, production, methods, domain, technology and release authority.	Mandates, named roles, concurrence points and unresolved authority questions.
Map critical assets and flows	Expose where meaning, quality, time or accountability is lost.	Stable identifiers, provenance, transformations, delays, controls and fallback.
Resolve priority dependencies	Avoid using AI to conceal weak metadata, access or coordination.	Owners, deadlines, resources and verification evidence.
Gate one AI assisted task	Test value without weakening official statistics.	Baseline, prompt record, benchmark, human review, security boundary and stop rules.

The institutional shift

An organization chart shows where units sit. An ecosystem map shows how information actually moves, where definitions change and where responsibility changes. That operational view is necessary before AI can be allocated safely.

AI enabled performance and retained authority

Module A1 is not satisfied when learners merely describe an ecosystem. They must use an approved AI capability to accelerate evidence work, verify the output independently and produce an accepted ecosystem map and decision record. The division of responsibility must be explicit.

Required field	Executive rule	Evidence
What AI could do	Extract candidate actors, assets, flows, definitions and inconsistencies from approved materials.	Prompt or configuration, source set and candidate output.
What AI should do	Perform only bounded tasks that are proportionate, lawful, secure and independently checkable.	Approved task allocation, restrictions and stop rules.
What the statistician must do	Define the statistical purpose, population, concepts, quality criteria, uncertainty and acceptance decision.	Named reviewer, verification record and signed disposition.
What the subject specialist must do	Validate contextual meaning, operating reality, domain conditions and consequences.	Evidence reviewed, concurrence, corrections or escalation.
How work is verified	Compare candidate records with authoritative sources, owners, registers and reproducible checks.	Claim ledger, correction log and residual limitations.
How performance improves	Compare baseline and AI assisted time under the same quality threshold.	Elapsed time, verification effort, rework and accepted output.

Non substitution boundary

AI may propose structures, summaries, classifications, questions and draft records. It may not establish official meaning, approve a source, determine fitness for use, accept residual risk or authorize release. It must never serve as the only source or the only verifier of its own output.

Priority risk picture

Risk	Why it matters	Immediate control
Conflicting concepts and revisions	Different figures can appear official while measuring different things.	Agree definitions, classifications, provenance and revision status.
Coverage and selection bias	Timely sources may omit critical populations or locations.	Measure coverage against the target population and validate by domain.
Missing identifiers	Linkage can create exclusions, false matches and untraceable transformations.	Govern linkage rules and quantify linkage error.
Narrow pilot evidence	Performance in accessible areas may not generalize nationally.	Use independent stratified validation and state the information boundary.
Fragile capacity	A pilot may depend on one analyst, temporary funding or a supplier.	Fund recurrent roles, reproducible pipelines, handover and fallback.
Confidentiality and unlawful reuse	Improper inputs can breach law, trust and statistical confidentiality.	Confirm legal basis, minimization, approved environment and disclosure review.

Kintaba decision example

The fictional Republic of Kintaba requires district crop production information six weeks before its annual survey estimate. Available evidence contains conflicting production totals, road related extension coverage bias, substantial weather station missingness, incompatible identifiers, geographically narrow satellite validation and fragile operating capacity.

Decision on the proposed automated estimate

Do not authorize an automated official estimate. First reconcile concepts and revisions, document the early estimate protocol, establish an independent benchmark, assess geographic and subgroup performance, confirm access and security, and name the human authorities for method, domain validation and release. A narrower pilot may use AI to organize discrepancy evidence or assist documentation, but generative AI must not become the estimator or release authority.

AI pilot authorization gate

Gate	Question requiring an evidenced answer	Required human authority
Purpose	What exact decision and statistical product will improve?	Sponsor and statistical lead approve intended use.
Baseline	What current timing, effort and accepted quality will be compared?	Process owner confirms comparability.
Data	Which sources are permitted, representative, documented and stable?	Data owner, steward and security authority concur.
Quality	Which benchmarks, domains, tolerances and uncertainty measures apply?	Statistician and methodologist approve criteria.
Domain meaning	Which contextual assumptions and consequences require specialist review?	Named subject matter specialist concurs or escalates.
Verification	Which checks are independent of the AI output?	Reviewer signs the verification and correction record.
Sustainability	Which recurrent skills, funding, infrastructure and fallback exist?	Service owner accepts recurrent responsibility.

Automatic stop conditions

- Protected information would enter an unapproved environment.
- The AI output would be its own sole benchmark or verifier.
- No qualified statistician or subject matter specialist can review a consequential claim.
- The proposed gain depends on weakening validation, concealing uncertainty or bypassing consultation.
- The system cannot reproduce the evidence trail or identify the version used.

Recommended ninety day leadership agenda

Period	Leadership action	Controlled output
Days 1 to 15	Appoint sponsor, statistical lead, domain lead, steward and security contact. Select one decision chain and record the baseline workflow.	Approved scope, roles, baseline and information classification.
Days 16 to 40	Map actors, authority, assets, flows, users, quality evidence and unknowns. Use AI only on approved materials and retain prompts.	Evidence backed map, claim ledger and prompt record.
Days 41 to 60	Verify every material claim; correct candidate records; prioritize no more than five gaps by consequence and dependency.	Accepted registers, correction log and owned action plan.
Days 61 to 75	Compare baseline and AI assisted effort under identical acceptance criteria. Screen one bounded use.	Quality adjusted time comparison and proceed, restrict or stop recommendation.
Days 76 to 90	If eligible, approve a controlled pilot protocol with independent review, stop rules and recurrent resources.	Signed pilot charter or documented decision not to pilot.

Measures of progress

- Critical actors, roles and authority are source traceable.
- Priority information assets have owners, provenance, access conditions, quality evidence and known changes.
- One critical flow is traceable from source to user decision, including transformations and delays.
- Material claims are classified as evidenced, reported or unknown.
- AI assisted work shows prompts, verification, corrections and quality adjusted time evidence.
- Every consequential output has named statistical and subject matter accountability.

Executive oversight dashboard

Dimension	Minimum indicator	Leadership interpretation
Statistical purpose	Decision, population, product and intended use approved.	No use case proceeds without a bounded public purpose.
Evidence integrity	Share of material claims independently verified and corrected.	Fluent AI output is not evidence.
Quality	Acceptance criteria met by baseline and AI assisted output.	Speed receives no credit below the quality threshold.
Efficiency	Net time saved after prompting, verification, correction and rework.	Measure quality adjusted time, not raw generation time.
Information protection	All inputs classified and processed in approved environments.	Any unauthorized exposure is a critical failure.
Human accountability	Named statistician and specialist sign off or escalate.	AI cannot approve its own work.
Transferability	Versioned prompt pattern, checklist or procedure with stopping rules.	A one off interaction is not institutional capability.

Executive decision record

Decision field	Record
Accountable sponsor	Name and position
Priority decision chain	Statistical product, user decision and deadline
Statistical lead	Name and position
Subject matter and stewardship leads	Names, positions and concurrence required
Security and legal contacts	Names and positions
Current AI authorization	Diagnostic and screening only; no automated release authority
Review date	Date for proceed, restrict, redesign or stop decision

Leadership position

The responsible position is neither to reject AI categorically nor to authorize it generically. Leadership should create a governed route from public purpose to evidence, from evidence to controlled testing, and from testing to an accountable decision. The central question is not whether a system can produce an output. It is whether the institution can explain what the output means, whom it represents, how uncertainty and error were assessed, who remains responsible and why the result is fit for the intended public use.

Proposed executive decision

Approve a ninety day national information ecosystem diagnostic for one priority statistical decision chain. Do not authorize automated official estimates or operational AI deployment through this decision. Require the diagnostic team to return with an evidence backed ecosystem map, a prioritized dependency register, a baseline and AI assisted comparison, a verification and correction record, and a recommendation to proceed, restrict, redesign or stop.

Mandatory conditions

Condition	Decision rule
AI assistance	Use only approved tools, approved information and bounded task specifications.
Professional authority	The statistician and relevant subject matter specialist retain non delegable responsibility.
Verification	No consequential claim is accepted without independent evidence or reproducible testing.
Efficiency	Count improvement only after verification burden, corrections and rework are included.
Release	No official figure, method or interpretation is released without authorized human approval.

Relationship to the Module A1 evidence package

This brief translates the canonical lesson into an executive decision. The learner workbook, facilitator guide, narrated presentation, executable notebook, data resources, exercises, assessment pack and institutional mapping toolkit provide the operational evidence. None substitutes for the others, and none transfers accountability from qualified humans to AI.